

Chapter 16 Segmentation

Base and bounds is not as flexible as we would like it to be:

- Requires big chunk of “free” space in the middle
- Taking up physical memory
- Doesn't look so bad with our small address spaces
 - Imagine a 32-bit address space (4 GB in size)
 - Typical program only use a few megabytes of memory

16.1 Segmentation: Generalized Base & Bounds

Idea: use multiple base/bounds pairs

- One base/bounds-pair for each **logical segment** of the address space

Def. Segment: *Contiguous position of address space of a particular length.*

Address space have three logically different segments

- Code
- Stack
- Heap




Figure 16.1: An Address Space (Again)

Example Fig 16.1: Each segment can be placed independently in physical memory.

Example Fig 16.2: 64 KB physical memory

Large amounts of unused address space can be accommodated.

HW: MMU support for segmentation

- set of three base and bounds register pairs

Segment	Base	Size
Code	32K	2K
Heap	34K	2K
Stack	28K	2K

Figure 16.3: Segment Register Values

Example: Translation (refer to Fig 16.1)

- Instruction fetch from virtual address 100 (code segment)
 - Translate: $100 + 32 \text{ KB} (32 \text{ KB} = 32 * 1024) = 32868$
- Heap: consider virtual address 4200 (heap segment)
 - Translate: $4200 - 4096 = 104 + 34 \text{ KB} = 34920$





Figure 16.1: An Address Space (Again)



16.2: Placing Segments In Physical Memory


16.2 Which Segment are we Referring to?

Q: How does the HW know which segment the address is referring to?

Explicit Approach:

- two top bits of the virtual address space — to indicate segment type

"00"	Code
"01"	Heap
"10"	Stack
"11"	Unused



From the Example above: Heap segment with virtual address 4200:

Segment bits: 01

Offset bits: 0b0000 0110 1000 = 0x068 = 104 decimal





Figure 16.1: An Address Space (Again)

Physical address = offset + base register

SEG_MASK = 0x3000 (bits 12 and 13: 11)

Example:

```
    01 0000 0110 1000
AND   11 0000 0000 0000
    01 0000 0000 0000
SHIFT 00 0000 0000 00 01
```

SEG_SHIFT = 12

In order to move the segment bits to the lowest two bits.

```
    01 0000 0110 1000
AND   00 1111 1111 1111
    00 0000 0110 1000
```

OFFSET_MASK = 0x0FFF (12 bits for the offset)

Base and Bounds variables are arrays with one entry per segment.

Pseudo code on the right can be used to compute the

Physical address from virtual address...

```
1  // get top 2 bits of 14-bit VA
2  Segment = (VirtualAddress & SEG_MASK) >> SEG_SHIFT
3  // now get offset
4  Offset  = VirtualAddress & OFFSET_MASK
5  if (Offset >= Bounds[Segment])
6      RaiseException(PROTECTION_FAULT)
7  else
8      PhysAddr = Base[Segment] + Offset
9      Register = AccessMemory(PhysAddr)
```

Implicit approach

- The hardware determine the segment based on how the virtual address was formed...
 - Generated from PC because of an instruction fetch — address is code segment
 - Generated from SP because the CPU is using it — address is stack segment
 - Any other address — heap segment

Another digression into binary and (simple) boolean algebra!

AND TABLE:

$a \text{ AND } b = c$

$0 \text{ AND } 0 = 0$

$0 \text{ AND } 1 = 0$

$1 \text{ AND } 0 = 0$

$1 \text{ AND } 1 = 1$

Example:

$x = 0001$

$y = 1101$

$z = 0001$

Bitwise AND: $x \& y = 0001$

Logical AND: $\text{len(array)} > 3 \&\& \text{timeSliceExpired} = \text{true/false}$

OR TABLE:

$a \text{ OR } b = c$

$0 \text{ OR } 0 = 0$

$0 \text{ OR } 1 = 1$

$1 \text{ OR } 0 = 1$

$1 \text{ OR } 1 = 1$

Example:

$x = 0001$

$y = 1101$

$z = 1101$

Bitwise OR: $a | b = 1101$

Logical OR: $\text{len(array)} > 3 || \text{timeSliceExpired} = \text{true/false}$

XOR TABLE: $a \oplus b = c$ $0 \oplus 0 = 0$ $0 \oplus 1 = 1$ $1 \oplus 0 = 1$ $1 \oplus 1 = 0$

16.3 What About the Stack?

Since the stack grows backwards — need extra bit in hardware to indicate the growing direction of the segment

Hardware can now translate such virtual address differently...

Example. Access virtual address 15KB (stack Fig 16.1).

15KB = 11 1100 0000 0000 (hex: 0x3C00)

Two top bits: 11 - stack segment

OFFSET_MASK: 0x0FFF

Offset: $0x0C00 = 3KB = 3072$

To find the correct negative offset (into the stack segment)

- Subtract max segment size (4 KB in this example) from 3 KB
- Negative offset: $3KB - 4KB = -1KB$
- Simply add negative offset to the base: $28KB + (-1KB) = 27 KB$ (physical address)
- (Maps to physical address 27KB in Fig 16.2)

Segment	Base	Size	Grows Positive?
Code	32K	2K	1
Heap	34K	2K	1
Stack	28K	2K	0

Figure 16.4: Segment Registers (With Negative-Growth Support)




Figure 16.1: An Address Space (Again)




Figure 16.2: Placing Segments In Physical Memory

16.4 Support for Sharing

Idea: save memory by sharing certain memory segments between address spaces (=process's)

HW: add protection bits

- A few extra bits per segment
- Read, Write, and Execute

Setting code segment to read-only:

- Multiple processes can share code
- Without harming isolation
 - Each process still thinks that it is accessing its own private memory
 - While OS is secretly sharing their memory!

HW: must check if a particular access is permissible

Example: if user process tries to

- Write to a read-only segment, or
- Execute from a non-executable segment,
- HW should raise an exception

Segment	Base	Size	Grows Positive?	Protection
Code	32K	2K	1	Read-Execute
Heap	34K	2K	1	Read-Write
Stack	28K	2K	0	Read-Write

Figure 16.5: Segment Register Values (with Protection)

16.6 OS Support

Motivation and Context

- Unused space between stack and heap need not be allocated in physical memory.
- Allowing for more address spaces (processes) in physical memory.

Tasks for the OS with segmentation:

- Q: What should the OS do on a context switch?
 - Segmentation registers must be saved and restored.
- Free space management
 - Find space for new processes
 - A process's address space has segments of different sizes.

General problem with free space management

- Physical memory becomes full of little holes of free space
- Makes it difficult to
 - Allocate new segments or
 - Grow existing ones
- This is called external fragmentation (Fig. 16.6)

Example Fig. 16.6: A process wish to allocate 20KB segment.

There is 24KB free, but not in one contiguous segment (rather in three non-contiguous chunks). OS cannot satisfy the 20KB request.

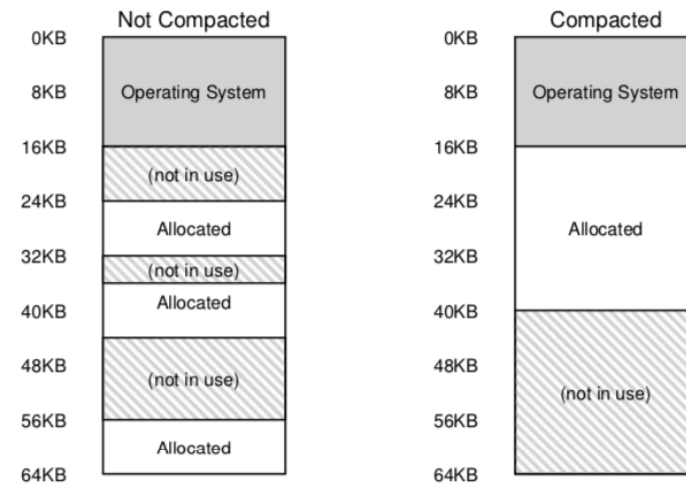


Figure 16.6: Non-compacted and Compacted Memory

One solution:

- Stop processes one-by-one, move segments, update segment registers to point to new location
- Expensive: memory-intensive and use of processor time

Simpler/Better approach:

- Use a free-list management algorithm
 - Keep large **extents** of memory available for allocation
- Ex. Best-fit algorithm — keeps list of free spaces — returns the one closest in size that satisfies the desired allocation for a memory allocation request.

16.7 Summary

Segmentation

- Support sparse address spaces
- Avoid wasting memory between logical segments of an address space
- Fast translation with low overhead
- Code sharing

Problems with segmentation (variable-sized segments)

- Free memory gets chopped into odd-sized pieces (external fragmentation)
- Memory allocation difficult (many smart algorithms exists)
 - But fundamentally hard to avoid external fragmentation
- Not flexible enough, to support fully generalized sparse address space
 - e.g., if heap is sparsely-used — must keep entire heap in memory ...