

Chapter 20 Paging: Smaller Tables

Problem with paging: various overheads

- Storage overhead for the page tables
- Page tables are stored in memory
- Requires extra memory lookup for each virtual address
 - Going to memory for translation information for every instruction fetch and for every load/store instruction
 - Prohibitively slow
- In chapter 19, we introduced the concept of an *address translation cache*, or TLB, to solve the latency problem.

Example. 32-bit address space (2^{32} bytes = 4 GB)

- 4 KB (2^{12} byte) pages
- 4 byte page-table entry

One address space has one million virtual pages ($2^{32} / 2^{12} = 2^{20}$).

Multiply this with the size of the page table entry: $4 \times 2^{20} = 4$ MB.

One page table ____

With 100 active ____

Q: How can we make page tables smaller?

- Avoid linear page tables (arrays) — they are too big
- Key idea: new data structures
- Challenge: other inefficiencies with new data structures

20.1 Simple Solution: Bigger Pages

Example with 32-bit address space again

But now let's consider 16 KB pages or 2^{14} bytes.

-

-

$2^{18} = 256k$ entries

Problem with approach:

-

-

Another alternative:

-

-

-

-

20.2 Hybrid Approach: Paging and Segments

Def. Hybrid: *two (or more) good ideas combined to achieve the best of both worlds.*

Example. Linear page table

-
-
-




Figure 20.1: A 16KB Address Space With 1KB Pages

Fig. 20.2 shows that most PTEs are invalid (unused)

Q: Why not have ____ ?

-

-

PFN	valid	prot	present	dirty
10	1	r-x	1	0
-	0	---	-	-
-	0	---	-	-
-	0	---	-	-
23	1	rw-	1	1
-	0	---	-	-
-	0	---	-	-
-	0	---	-	-
-	0	---	-	-
-	0	---	-	-
-	0	---	-	-
-	0	---	-	-
-	0	---	-	-
-	0	---	-	-
-	0	---	-	-
28	1	rw-	1	1
4	1	rw-	1	1


Figure 20.2: A Page Table For 16KB Address Space

Use base and bounds, but now:

- **Base** points to physical address of the page table for the segment
 - Code, stack, heap
- **Bounds** indicate end of page table
 - How many valid pages it has

Example. 32-bit Virtual Address Space

- 2 bit segment
- 18 bits for VPN
- 12 bits for offset (4 KB pages)
- HW: 3 base/bounds pairs of registers



On context switch:

-

On TLB miss:

- HW
-
- HW takes
-
-

```
SN          = (VirtualAddress & SEG_MASK) >> SN_SHIFT
VPN         = (VirtualAddress & VPN_MASK) >> VPN_SHIFT
AddressOfPTE = Base[SN] + (VPN * sizeof(PTE))
```

Only difference from linear page tables:

-

-

Example. If *code* segment consumes three pages (0, 1, 2):

- Page table will have only three entries
- Bounds register will be set to 3
 - Memory accesses beyond 3 leads to exception — process termination

Benefits with hybrid approach:

-
-
-

Challenges with hybrid approach:

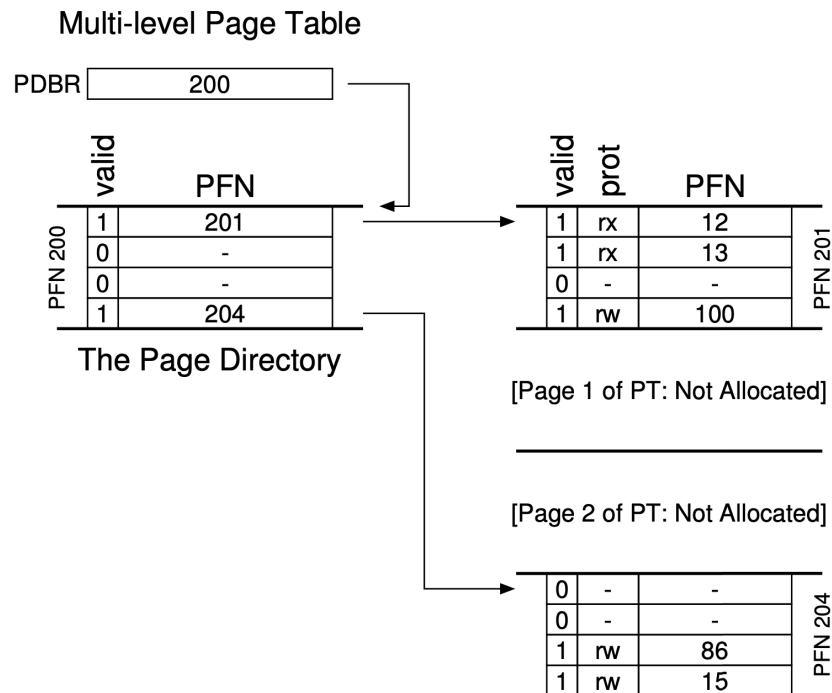
-
-
-
-
-

20.3 Multi-level Page Tables

Tree of linear page tables.

Basic idea:

- Chop Linear Page Table into page-sized units
- If an entire page of PTEs is invalid — don't allocate page at all
- **Page directory**: data structure to track
 - If a page in PT is valid and
 - Location in memory if valid

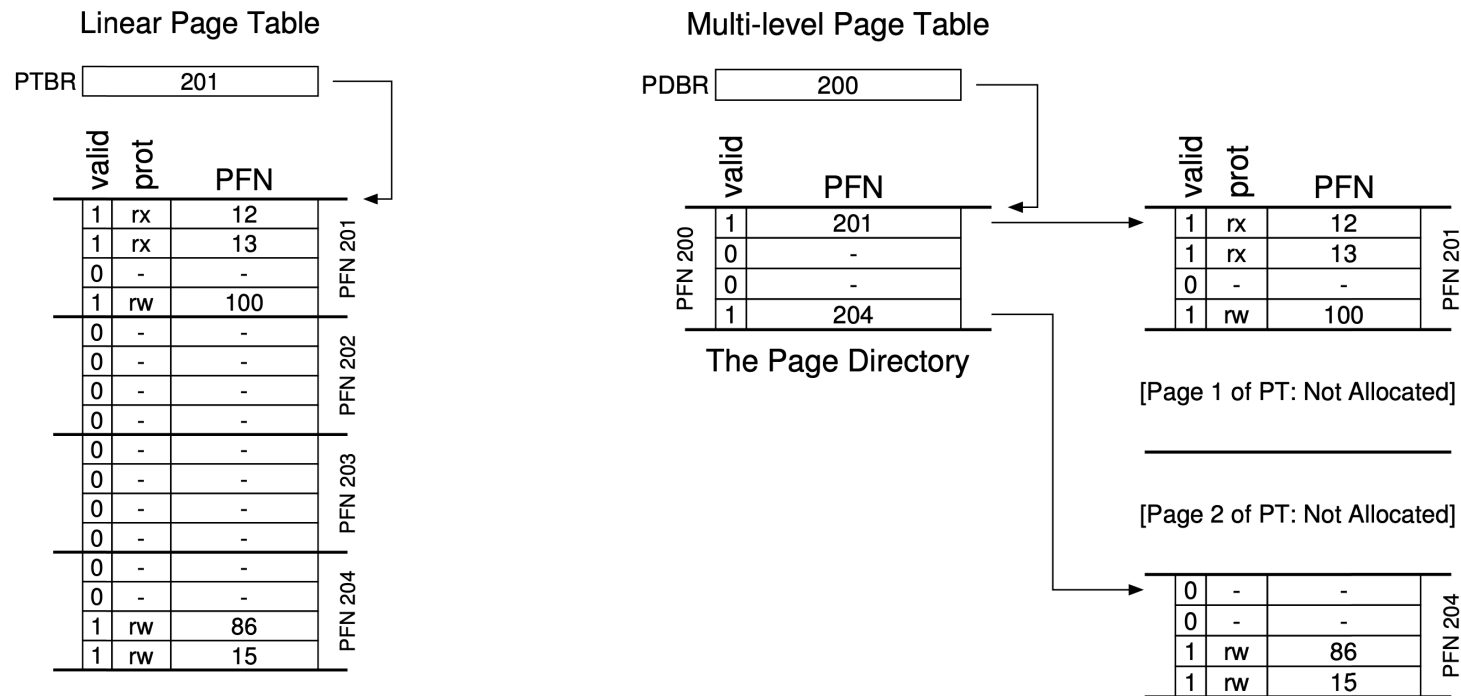


Linear Page Table:

- Two middle pages (PFN202 and PFN203) are not used
 - They have no mappings... but they still consume space in the linear PT.

Multi-level Page Table:

- Page Directory points to PFNs containing PFNs.
 - There are no pages in the middle two entries in the PD. Hence not allocated.



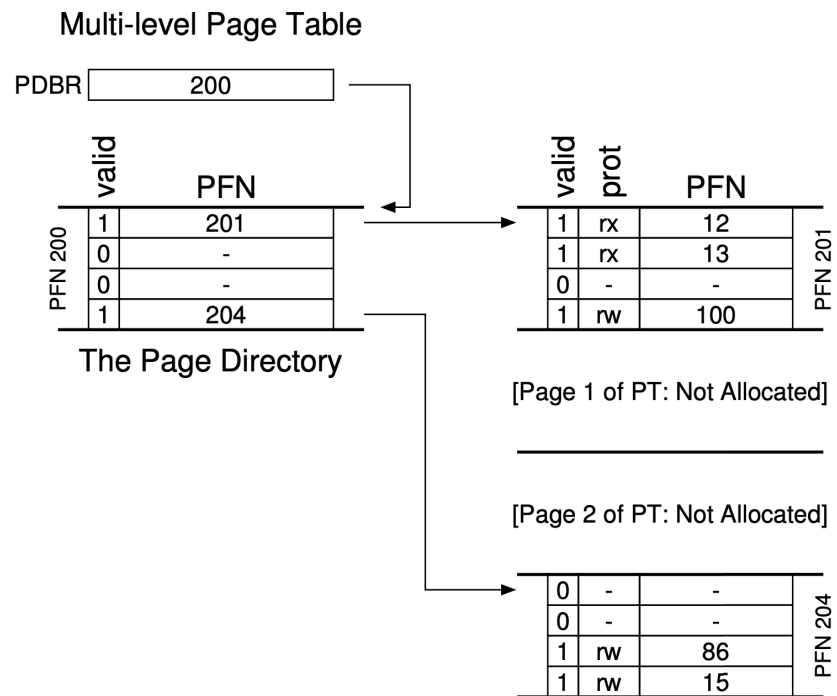
Query Page Directory to find out

- Where a page of the PT is
- Or the entire page of the PT contains no valid pages

Page Directory contains **Page Directory Entries** (PDEs) — similar to PTE.

Valid bit in the PDE:

- 1 means at least one PTE exists for page
- 0 means none of the PTEs in the PDEs range is valid



Advantages of Multi-level Page Tables

- Allocate only PT space for the addresses being used
 - Compact and supports sparse address spaces
- Page of PT fits on a page
 - Easier to manage memory
 - OS: Get next free page to grow the PT
 - (Not possible to grow a linear page table)
- We add a **level of indirection** by use of Page Directory.
 - Allow placing PT pages wherever in physical memory
- On TLB hit, fast, no performance penalty.

Drawbacks with Multi-level Page Tables

- On TLB miss:
 - Two loads from memory required to get translation
 - One for the PD and one for the PTE itself
- Linear Page table: just one load from memory
- Complexity: more involved than simple Linear PT lookup!

Time-space tradeoff: give some time to save some space

We make PT lookups more complicated to save space.

Example. Utilization in Linear Page Table

Small AS (16 KB), 64-byte pages

14-bit virtual AS:

- 8 bits for VPN
- 6 bits for offset
- Linear PT: $2^8 = 256$ entries
 - VP 0, 1: code
 - VP 4, 5: heap
 - VP 254, 255: stack
 - Rest unused

Utilization = $(2+2+2)/256 = 2,3 \%$

0000 0000	code
0000 0001	code
0000 0010	(free)
0000 0011	(free)
0000 0100	heap
0000 0101	heap
0000 0110	(free)
0000 0111	(free)
.....	... all free ...
1111 1100	(free)
1111 1101	(free)
1111 1110	stack
1111 1111	stack


Figure 20.4: A 16KB Address Space With 64-byte Pages

Example. Build two-level Page Table

- Linear PT: 256 entries, PTE = 4 bytes
 - PT size = $256 * 4$ bytes = 1 KB
- Break Linear PT up into page-sized units: 64 bytes
 - 1 KB Page Table divided into 16 64-byte pages
 - Each page can hold $64 / 4 = 16$ PTEs

Page-Directory Index (PDIndex)

- 256 entries, spread over 16 pages
- PD: need one entry per page: 16 entries (or 4 bits of the VPN)
- PDIndex is used to compute address of the PDE
 - $PDEAddr = PageDirBase + (PDIndex * sizeof(PDE))$




- Load PDE from the PDEAddr
 - If PDE invalid: raise exception
 - Else PDE valid

Next: Load PTE from PT pointed to by PDE

Page-Table Index (PTIndex)

- Index into PT itself
 - Gives us the address of our PTE
- PTIndex is the remaining 4 bits of the VPN
 - $PTEAddr = (PDE.PFN \ll SHIFT) + (PTIndex * sizeof(PTE))$
 - Note: the PDE's PFN must be left shifted into place before combining with the PTIndex



Example. Multi-level Page Table in Fig. 20.5

Page Directory (left)

- Use only two regions of address space
 - (2 of 16 entries)
- In-between space is unused

Page Table for PFN 100 (middle)

- Use four pages (0, 1, 4, 5) of address space
 - (4 of 16 entries)
- Code pages: 0, 1 ; Heap pages: 4, 5

Page Directory		Page of PT (@PFN:100)			Page of PT (@PFN:101)		
PFN	valid?	PFN	valid	prot	PFN	valid	prot
100	1	10	1	r-x	—	0	—
—	0	23	1	r-x	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	80	1	rw-	—	0	—
—	0	59	1	rw-	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	55	1	rw-
101	1	—	0	—	45	1	rw-

Figure 20.5: A Page Directory, And Pieces Of Page Table

Page Table for PFN 101 (right)

- Use two pages (14, 15) of address space
 - (2 of 16 entries)
- Stack pages: 254, 255

Utilization: $2+4+2 / (3*16) = 8 / 48 = 16,7 \%$ (of the three pages)

But more important metric:

We now only use 3 pages instead of 16 pages

for storing the PT: $(16-3) / 16 = 81 \%$ improvement

Savings for large addresses spaces are much greater!

Translation example:

Virtual address 0x3F80 = 1111_1110_000000

- PDIndex = 1111 (15 th in Page Directory) —> 101
- PTIndex = 1110 (14 th entry in Page of PT (@PFN 101))
 - -> PFN = 55 = 0011_0111
 - Offset = 000000
 - SHIFT = 6 (since offset is 6 bits)

PhysAddr = (PTE.PFN << SHIFT) + Offset
 = 0011_0111_000000
 = 0x0DC0

Page Directory		Page of PT (@PFN:100)			Page of PT (@PFN:101)		
PFN	valid?	PFN	valid	prot	PFN	valid	prot
100	1	10	1	r-x	—	0	—
—	0	23	1	r-x	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	80	1	rw-	—	0	—
—	0	59	1	rw-	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	55	1	rw-
101	1	—	0	—	45	1	rw-


Figure 20.5: A Page Directory, And Pieces Of Page Table

More than Two Levels

For large address spaces, we need deeper multi-level page tables.

Idea (not completely accurate):

- Lookup using PD Index 0 (high order bits) to find *index* into second-level PD
- Combine this *index* with PD Index 1 to find *index2* into PT
- Combine this *index2* with PT Index and offset to find physical address



The Translation Process (now with TLB)

TLB hit: all good, no performance penalty for multi-level page table lookup

TLB miss: HW must do full multi-level lookup

- For k -level page table:
 - We get k additional memory accesses versus just one to get the actual data/instruction we are interested in.

```
1  VPN = (VirtualAddress & VPN_MASK) >> SHIFT
2  (Success, TlbEntry) = TLB_Lookup(VPN)
3  if (Success == True)    // TLB Hit
4      if (CanAccess(TlbEntry.ProtectBits) == True)
5          Offset = VirtualAddress & OFFSET_MASK
6          PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
7          Register = AccessMemory (PhysAddr)
8      else
9          RaiseException(PROTECTION_FAULT)
10 else    // TLB Miss
11     // first, get page directory entry
12     PDIndex = (VPN & PD_MASK) >> PD_SHIFT
13     PDEAddr = PDBR + (PDIndex * sizeof(PDE))
14     PDE = AccessMemory (PDEAddr)
15     if (PDE.Valid == False)
16         RaiseException(SEGMENTATION_FAULT)
17     else
18         // PDE is valid: now fetch PTE from page table
19         PTIndex = (VPN & PT_MASK) >> PT_SHIFT
20         PTEAddr = (PDE.PFN << SHIFT) + (PTIndex * sizeof(PTE))
21         PTE = AccessMemory (PTEAddr)
22         if (PTE.Valid == False)
23             RaiseException(SEGMENTATION_FAULT)
24         else if (CanAccess(PTE.ProtectBits) == False)
25             RaiseException(PROTECTION_FAULT)
26         else
27             TLB_Insert(VPN, PTE.PFN, PTE.ProtectBits)
28             RetryInstruction()
```

Figure 20.6: Multi-level Page Table Control Flow

20.6 Summary

Tradeoffs:

- Time vs space: Sacrifice time to save space
- Complexity: Make PT lookups more complicated to save space for PTs