

# Chapter 28 Locks

Introduce lock synchronization primitive to protect critical sections.

```
mu = sync.Mutex{}
```

```
mu.Lock()
```

```
counter = counter + 1 // code in the critical section
```

```
mu.Unlock()
```

## **Declare a lock variable: mutex (mutual exclusion)**

- Holds the state of the lock at any instant in time
  - Available — unlocked and free
    - No thread holds the lock
  - Acquired — locked or held
    - Exactly one thread holds the lock
- Could also hold other info in the lock/mutex data structure
  - Which thread holds the lock
  - A queue for ordering lock acquisitions

### **Semantics of Lock() and Unlock():**

- Lock():
  - The first thread will acquire the lock (and can enter the CS)
  - Becomes owner of lock
  - Another thread calling Lock() on the same mutex
    - Will not return (block): thread is waiting
    - While other thread is in the CS
- Unlock():
  - Owner of mutex calls Unlock() (when leaving the CS)
    - Lock becomes free
  - If no other threads are waiting for the lock
    - State of lock is set to free
  - Otherwise: there are threads waiting stuck in the Lock() method
    - One of them will acquire the lock (and can enter the CS)

## **Lock Granularity**

- Coarse-grained locking
  - One “big” lock
    - Covering large/entire data structure
- Fine-grained locking
  - Several locks
    - Each lock covering different positions of a data structure
  - Benefit: increased concurrency compared to a single “big” lock for all CSs

## Thread API Guidelines (from Chapter 27)

- Keep it simple
  - Any code to lock or signal between threads should be *as simple as possible*
  - Tricky thread interactions lead to subtle bugs (deadlocks)
  - Applies to Go's channel-based interactions as well
- Minimize thread interactions
  - Keep the number of ways in which threads interact to a minimum
  - Each interaction should be carefully thought out and constructed with well-known patterns
- Initialize locks and condition variables (in C)
  - Failure to initialize: sometimes works fine and sometimes fails in strange ways

## Thread API Guidelines (continued)

- Check return codes:
  - Return codes in C and Unix contain info: should be checked
- Be careful about how you pass arguments to and return values from threads
  - If pass by reference to a variable allocated on the stack: you are doing it wrong!
- Each thread has its own stack
  - A thread's locally allocated variables should be considered private
  - To share data between threads
    - Allocate space on the heap
    - Or use a global variable
- Always use condition variables (CVs) to signal between threads
  - Don't use a simple flag!
  - Good alternative to CVs: Channels in Go.

## 28.4 Evaluating Locks

Evaluation criteria

- **Basic task/Correctness:** does it provide mutual exclusion?
- **Fairness:** does each thread contenting for the lock get a fair shot at acquiring the lock once it is free?
- **Performance:** time overheads added by using a lock. Cases to consider:
  - No contention
  - Multiple threads contending for the lock on
    - A single CPU
    - Multiple CPUs

## 28.5 Controlling Interrupts

Early solutions for single CPU systems:

```
func Lock() {                               func Unlock() {  
    DisableInterrupts()                     EnableInterrupts()  
}
```

HW instruction to turn off interrupts  
And re-enable interrupts again (unlock)

Pro/Con:

- + Easy to understand that no other thread or the OS can interrupt during the CS
- Requires calling thread to perform *privileged* instruction
  - Must trust thread / arbitrary program
  - Greedy program: call Lock() at beginning of execution and monopolize the CPU
  - Malicious program: call Lock() and enter infinite loop...
  - Only recourse: restart system
- Does not work on multiprocessor systems
  - Each CPU has their own interrupts
- Disabling interrupts for long time: lead to lost interrupts
  - CPU may miss I/O completion event
- Slow to disable / enable interrupts

## 28.6 Lock Implementation Using Loads/Stores

Let's use a single **flag** variable

- Access flag variable via normal memory load/store
- Will it be sufficient?

If a thread calls lock() when it is being held by another thread, it will spin-wait until the thread calls unlock()

```
1  typedef struct __lock_t { int flag; } lock_t;
2
3  void init(lock_t *mutex) {
4      // 0 -> lock is available, 1 -> held
5      mutex->flag = 0;
6  }
7
8  void lock(lock_t *mutex) {
9      while (mutex->flag == 1) // TEST the flag
10         ; // spin-wait (do nothing)
11     mutex->flag = 1;         // now SET it!
12 }
13
14 void unlock(lock_t *mutex) {
15     mutex->flag = 0;
16 }
```

Figure 28.1: First Attempt: A Simple Flag

Thread 1	Thread 2
call lock()	
while (flag == 1)	
interrupt: switch to Thread 2	
	call lock()
	while (flag == 1)
	flag = 1;
	interrupt: switch to Thread 1
flag = 1; // set flag to 1 (too!)	

Figure 28.2: Trace: No Mutual Exclusion

Two problems:

- Correctness
- Performance

Fig. 28.2 Shows an interleaving that both threads get the lock — and can enter the CS.

*This means that we have no **mutual exclusion**, violating the correctness.*

Performance: **spin-waiting** wastes time and CPU cycles waiting for another thread to release the lock

- On uniprocessor: the thread that the waiter is waiting for can't even run until a timer interrupt.



## 28.7 Spin Locks with Test-and-Set

HW support

- test-and-set instruction
- Intel x86: **xchg** instruction

The TestAndSet() func is performed atomically.

- Copies and returns the old value
- Update ptr with a provided new value

```
1      int TestAndSet(int *old_ptr, int new) {
2          int old = *old_ptr; // fetch old value at old_ptr
3          *old_ptr = new;     // store 'new' into old_ptr
4          return old;         // return the old value
5      }
```

This TAS instruction makes it possible to implement a simple **spin lock**.

Allows to test the old value (returned) while setting the memory to the new value.

```
1  typedef struct __lock_t {
2      int flag;
3  } lock_t;
4
5  void init(lock_t *lock) {
6      // 0 indicates that lock is available, 1 that it is held
7      lock->flag = 0;
8  }
9
10 void lock(lock_t *lock) {
11     while (TestAndSet(&lock->flag, 1) == 1)
12         ; // spin-wait (do nothing)
13 }
14
15 void unlock(lock_t *lock) {
16     lock->flag = 0;
17 }
```

Figure 28.3: A Simple Spin Lock Using Test-and-set

First case:

- A thread calls lock() when no other threads hold the lock
  - Flag is 0
- When thread calls TAS(flag, 1)
  - TAS returns old value of flag, which is 0
- Thread will not spin since  $\text{TAS}(\text{flag}, 1) = 0 \neq 1$
- TAS will set flag to 1 (atomically)
  - Indicating that lock is now held
- On exiting CS and call unlock()
  - The thread sets flag back to 0.

```

1  typedef struct __lock_t {
2      int flag;
3  } lock_t;
4
5  void init(lock_t *lock) {
6      // 0 indicates that lock is available, 1 that it is held
7      lock->flag = 0;
8  }
9
10 void lock(lock_t *lock) {
11     while (TestAndSet(&lock->flag, 1) == 1)
12         ; // spin-wait (do nothing)
13 }
14
15 void unlock(lock_t *lock) {
16     lock->flag = 0;
17 }

```

Figure 28.3: A Simple Spin Lock Using Test-and-set

Second case:

- A thread calls lock() when another thread hold the lock
  - Flag is 1
- When this thread calls TAS(flag, 1)
  - It returns the old value of flag which is 1
  - Hence the thread enters spin-waiting
    - Repeatedly checking the flag using the TAS instruction
- Only when another thread sets the flag to 0
  - Will “this thread” call TAS() again, and hopefully returning 0.
  - While atomically setting flag to 1, and acquiring the lock and can enter the CS

Simplest type of lock: Spin Lock

Simply spins, wasting CPU cycles until lock becomes available.

To work on a single CPU, it requires a **preemptive scheduler**

- Interrupt threads periodically to run OS scheduler and replace threads...

## 28.8 Evaluating Spin Locks

Correctness: YES!

Fairness:

- Simple spin locks are not fair and may lead to starvation
- A thread may spin forever under contention
  - Other threads may grab the lock in front of a spinning thread

Performance:

- Single CPU case: Overheads can be quite “painful”
  - Thread holding lock preempted in CS
  - Scheduler runs all other threads
    - Each tries to acquire lock
    - Each thread spin for the duration of a time slice (until giving up)
    - Wasted an entire time slot

Performance:

- Multiple CPUs case: Can work reasonably well
  - Thread A on CPU 1
  - Thread B on CPU 2
  - A and B contend for the lock
    - If A gets the lock
    - B tries to get the lock, B will spin (on CPU2)
    - Assuming the CS is short
    - Lock quickly becomes available
    - And gets acquired by B

*Spinning to wait for lock held by another processor can be effective.*

## 28.9 Compare-And-Swap

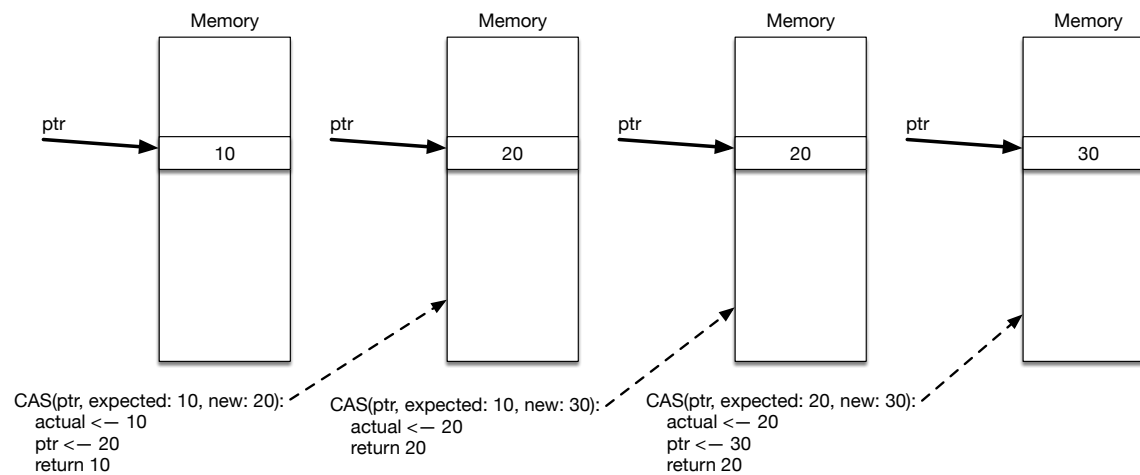
Slightly more powerful primitive than Test-And-Set.

Not needed for mutex locks.

For lock-free synchronization, we need CAS.

```
1  int CompareAndSwap(int *ptr, int expected, int new) {  
2      int actual = *ptr;  
3      if (actual == expected)  
4          *ptr = new;  
5      return actual;  
6  }
```

Figure 28.4: Compare-and-swap



### 28.13 A Simple Approach: Just Yield, Baby

First attempt:

- When you are going to spin
  - Instead: give up CPU to another thread
- Assume OS primitive: **yield()**

```
1  void init() {
2      flag = 0;
3  }
4
5  void lock() {
6      while (TestAndSet(&flag, 1) == 1)
7          yield(); // give up the CPU
8  }
9
10 void unlock() {
11     flag = 0;
12 }
```

Recall: Thread state: **running**, **ready**, or **blocked**.

When a thread calls `yield()`, it deschedules itself.

Moving from **running** to **ready** state.

**Example Two threads;** A holds the lock; B want to get the lock.

Instead of B spinning, B simply yields the CPU, allowing A to run again and finish its CS.

**Example.** 100 threads contenting for a lock, repeatedly.

One thread acquires the lock — and is preempted before finishing its CS (and thus not releasing the lock)

The other 99 threads call `Lock()`, find that it is held, and yield the CPU.

99 threads will run-and-yield before the thread holding the lock gets to run again.

Costly context switches (for no good reason).



## 28.14 Using Queues: Sleeping Instead of Spinning

Exert more control over which thread gets to acquire the lock next

- OS support
- Queue to keep track of threads waiting to acquire the lock

```
1  typedef struct __lock_t {
2      int flag;
3      int guard;
4      queue_t *q;
5  } lock_t;
6
7  void lock_init(lock_t *m) {
8      m->flag = 0;
9      m->guard = 0;
10     queue_init(m->q);
11 }
12
13 void lock(lock_t *m) {
14     while (TestAndSet(&m->guard, 1) == 1)
15         ; //acquire guard lock by spinning
16     if (m->flag == 0) {
17         m->flag = 1; // lock is acquired
18         m->guard = 0;
19     } else {
20         queue_add(m->q, gettid());
21         m->guard = 0;
22         park();
23     }
24 }
25
26 void unlock(lock_t *m) {
27     while (TestAndSet(&m->guard, 1) == 1)
28         ; //acquire guard lock by spinning
29     if (queue_empty(m->q))
30         m->flag = 0; // let go of lock; no one wants it
31     else
32         unpark(queue_remove(m->q)); // hold lock (for next thread!)
33     m->guard = 0;
34 }
```

Figure 28.9: Lock With Queues, Test-and-set, Yield, And Wakeup

Use two locks

- A **guard** lock to protect the lock\_t data structure (queue and flag)
  - Lock() and Unlock() must acquire the **guard** lock
  - Before updating the queue and flag variables
- The lock itself is the flag variable
- The **guard** lock is a spinning lock
  - But the time spent spinning is limited to the few instructions in the lock() and unlock() code
- Much better than scenarios where user-defined and arbitrarily long CSs.

If flag = 0: we get the lock

Otherwise:

- Add ourselves to the queue
- Release the guard lock
- Park ourselves

Unlock():

- Acquire the guard lock
- If queue empty: release lock
- Else:
  - Pass lock onto the next thread waiting in queue: Unpark()
  - When the “unparked” thread is woken up by the scheduler, it will resume after the park(), line 23
- Release the guard lock

