

Chapter 31 Semaphores

Single primitive for all things related to synchronization: can replace both

- Locks and
- Condition variables

Additional use of semaphores:

- Worker pool
 - Sometimes not desirable to create thousands or millions of goroutines
 - Instead we can limit the number of goroutines that can run at the same time
 - To some reasonable number, e.g., the number of CPU cores
 - Use a pool of “worker” goroutines/threads
 - Schedule work on these “workers”

Def. Semaphore: An object with an integer value that can be manipulated by

- `sem.Wait() = P()`
- `sem.Post() = V()`

Before using a semaphore:

- Initialize with some value
- `sem := &Semaphore{value: 1}`

Behavior of semaphore functions:

```
func (s *Semaphore) Wait() int {  
    decrements the value of semaphore s by one  
    wait if value of s is negative  
}
```

```
func (s *Semaphore) Post() int {  
    increments the value of semaphore s by one  
    if there are one or more threads waiting, wake one  
}
```

- Wait():
 - Returns right away if $s.value > 1$
 - Otherwise, suspends execution, waiting for a subsequent Post() call
 - Multiple calling threads: queued waiting to be woken up
- Post():
 - Increments $s.value$
 - If s has waiting threads, wake one up

Invariant of a semaphore object:

- $s.value < 0$ implies that $s.value == \# \text{waiting threads}$

31.2 Binary Semaphore (Locks)

Use a semaphore as a lock:

```
sem := &Semaphore{value: X} // initialize semaphore to X
sem.Wait();                  // Lock()
// critical section
sem.Post();                   // Unlock()
```

Q: What should X be?

X=1

Case 1: Thread 0 acquires and releases lock without Thread 1 interfering.

Value of Semaphore	Thread 0	Thread 1
1		
1	call sem_wait()	
0	sem_wait() returns	
0	(crit sect)	
0	call sem_post()	
1	sem_post() returns	

Figure 31.4: Thread Trace: Single Thread Using A Semaphore

Case 2: Thread 0 holds the lock, when Thread 1 tries to enter critical section

Value	Thread 0	State	Thread 1	State
1		Running		Ready
1	call sem_wait()	Running		Ready
0	sem_wait() returns	Running		Ready
0	(crit sect: begin)	Running		Ready
0	Interrupt; Switch→T1	Ready		Running
0		Ready	call sem_wait()	Running
-1		Ready	decrement sem	Running
-1		Ready	(sem<0)→sleep	Sleeping
-1		Running	Switch→T0	Sleeping
-1	(crit sect: end)	Running		Sleeping
-1	call sem_post()	Running		Sleeping
0	increment sem	Running		Sleeping
0	wake(T1)	Running		Ready
0	sem_post() returns	Running		Ready
0	Interrupt; Switch→T1	Ready		Running
0		Ready	sem_wait() returns	Running
0		Ready	(crit sect)	Running
0		Ready	call sem_post()	Running
1		Ready	sem_post() returns	Running

Figure 31.5: Thread Trace: Two Threads Using A Semaphore

31.3 Semaphores for Ordering

Semaphores can also be used to order events in a concurrent program

- Usage pattern
 - One thread waiting for something to happen
 - Another thread making that something happen
 - Signaling that it happened
 - Waking the waiting thread
- Similar to use of condition variables

Case 1: Parent thread continues to run and reach sem_wait() first

Value	Parent	State	Child	State
0	create (Child)	Running	(Child exists; is runnable)	Ready
0	call sem_wait ()	Running		Ready
-1	decrement sem	Running		Ready
-1	(sem<0) →sleep	Sleeping		Ready
-1	Switch→Child	Sleeping	child runs	Running
-1		Sleeping	call sem_post ()	Running
0		Sleeping	increment sem	Running
0		Ready	wake (Parent)	Running
0		Ready	sem_post () returns	Running
0		Ready	Interrupt; Switch→Parent	Ready
0	sem_wait () returns	Running		Ready

Figure 31.7: Thread Trace: Parent Waiting For Child (Case 1)

Case 2: Child thread runs first reach sem_post() before parent runs sem_wait()

Value	Parent	State	Child	State
0	create (Child)	Running	(Child exists; is runnable)	Ready
0	Interrupt; Switch→Child	Ready	child runs	Running
0		Ready	call sem_post ()	Running
1		Ready	increment sem	Running
1		Ready	wake (nobody)	Running
1		Ready	sem_post () returns	Running
1	parent runs	Running	Interrupt; Switch→Parent	Ready
1	call sem_wait ()	Running		Ready
0	decrement sem	Running		Ready
0	(sem ≥ 0) → awake	Running		Ready
0	sem_wait () returns	Running		Ready

Figure 31.8: Thread Trace: Parent Waiting For Child (Case 2)

31.4 The Producer/Consumer Problem (Bounded Buffer)

First Attempt: Use two semaphores

- Empty
- Full

MAX=1: This works even with multiple producers and multiple consumers

Next, let's assume multiple producers and multiple consumers when

MAX=10

Q: What can happen if MAX=10?

A: There is nothing that protects the critical section around `get()` and `put()` so multiple threads can enter CS.

That is, `sem.Wait()` only waits when negative, and does not prevent another thread from entering the CS.

Example: T1 and T2, both calling put() concurrently.

- Assume T1 runs first, and fill buffer (fillIndex = 0 on Line F1 in the Go code)
- T1 descheduled before Line F2
- T2 runs, and puts its data at the 0th element
 - Overwrite the data just written by T1
 - We have data loss!

Solution: Add mutual exclusion!

- Guard calls to put() and get() with locks (binary semaphore)
- However, our solution will cause a deadlock... Why?

Example: Illustrate when our code can deadlock (Fig. 31.11)

- C runs first
- Gets lock: `mutex.Wait()` ; Line c0
- Calls `full.Wait()` ; Line c1 — but there is nothing to consume
 - Yield the CPU, but still holds the lock
- Next, P runs to produce some data
 - Tries to get lock: `mutex.Wait()` ; Line p0
 - Lock already held; blocks!
- Simple cycle:
 - C holds *mutex*, waiting for someone to **signal** on the *full* semaphore
 - P could **signal** on *full* semaphore, but is **waiting** for the *mutex*.
- Both C and P are stuck waiting for each other!
 - Classic deadlock!

31.5 Reader-Writer Locks

Distinguish between operations that

- Read (often much more frequent)
- Write

to a data structure.

Example: Concurrent list operations:

- Insert into list — update the list structure (WRITE)
- Lookup in the list — doesn't change the list structures (READ)

Since lookups don't change the list structure

- Can have many lookups run concurrently without causing problems since they are just reading that data
- Works if we can ensure that no writers run concurrently

This is the job of a reader-writer lock.

Pseudo-code for ReadWrite Lock:

```
func (rw *RWLock) AcquireReadLock() {  
    wait for rw.lock: lock to access internal rwlock data structure  
    readers++  
    if readers == 1  
        // first reader  
        wait for the write lock  
    release the rw.lock  
}
```

```
func (rw *RWLock) ReleaseReadLock() {  
    behavior is similar to acquire read lock  
    except it releases the write lock when all readers are done  
}
```

Discussion:

- First reader to get the writelock, essentially allows any reader to get the readlock too.
- A writer must thus wait for all readers to finish!

Problem with this implementation: Fairness

- Easy for readers to starve writers

A possible fix:

- Prevent more readers from entering lock once a writer is waiting

31.6 The Dining Philosophers

Intellectually stimulating / interesting:

- But practical utility is low
- Included because everyone should know about the problem

Setup:

- Five philosophers sitting around a table
- Single chopstick between each philosopher




Figure 31.14: The Dining Philosophers

The philosophers alternate between

- Eating: need two chopsticks
 - One on its right and one on its left
- Thinking: don't need any chopsticks

```
while true {  
    think()  
    getforks()  
    eat()  
    putforks()  
}
```




Figure 31.14: The Dining Philosophers

- Write `getforks()` and `putforks()`

Requirements:

- No deadlock
- No philosopher starves (never gets to eat)
- Concurrency is high (as many philosophers can eat at the same time as possible)

Helper functions:

```
int left(int p) {    return p;    }  
int right(int p) {   return (p+1) % 5;   }
```

Explain:

- left(): refers to the chopstick to the left of the philosopher p .
- right(): similar. Modulo operator allow last philosopher $p=4$ to get its right chopstick, which is 0.




Figure 31.14: **The Dining Philosophers**

We will use five semaphores; one for each chopstick

```
func getforks() {  
    sem.Wait(forks[left(p)]);  
    sem.Wait(forks[right(p)]);  
}
```

```
func putforks() {  
    sem.Post(forks[left(p)]);  
    sem.Post(forks[right(p)]);  
}
```

First: pick up the left chopstick and then pick up the right chopstick.

Deadlock: *If each philosopher grab the fork on their left before any philosopher can grab the fork on their right, each will be stuck with one fork, waiting for another, forever.*

p0, f0, p1, f1, p2, f2, p3, f3, p4, f4 ... deadlock.

All forks acquired.

All philosophers stuck waiting for another fork.

A Solution: Breaking the Dependency

Make one of the philosophers pick up the chopsticks in a different order.

- Right, then left instead of
- Left the right.

Breaks the cycle of waiting.

```
void getforks() {  
    if p == 4 {  
        sem.Wait(forks[right(p)]);  
        sem.Wait(forks[left(p)]);  
    } else {  
        sem.Wait(forks[left(p)]);  
        sem.Wait(forks[right(p)]);  
    }  
}
```




Figure 31.14: The Dining Philosophers

p0, f0, p1, f1, p2, f2, p3, f3, p4, f0 (must wait), p3, f4, ... not deadlock

31.8 Summary

Semaphores can be viewed as a generalization of locks and condition variables

- Turns out building CVs using semaphores is difficult!
- Yet, some programmers use semaphores exclusively, because of their “simplicity” and utility.
- Arguably locks are simpler!!
- Condition variables are still a bit complex!
- Go channels are often much easier than CVs.
- Semaphores are useful for work load management!