

I/O Devices (Chapter 36)

How should I/O be integrated into systems?

What are the general mechanisms?

How can we make them efficient?

I/O Devices (Chapter 36)

36.1 System Architecture (Generic prototype)

PCI: Peripheral Component Interconnect

SCSI: Small computer System Interface

SATA: Serial Advanced Technology Attachment

USB: Universal Serial Bus

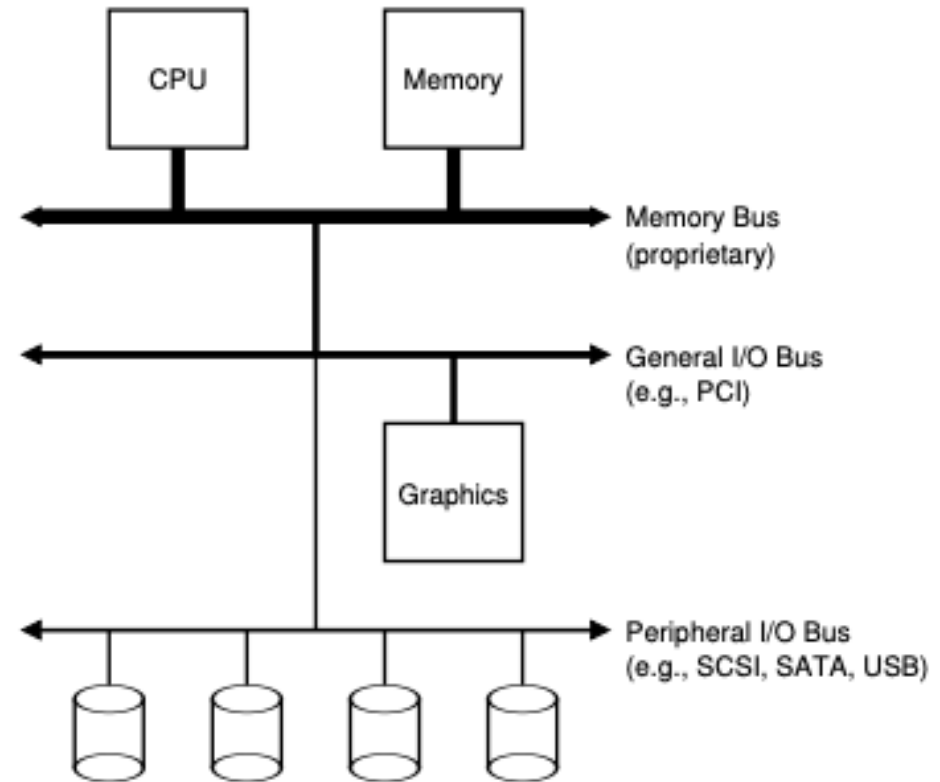


Figure 36.1: Prototypical System Architecture

I/O Devices (Chapter 36)

36.1 System Architecture Intel's Z270 Chipset

PCIe: Peripheral Component Interconnect Express

eSATA: Express Serial Advanced Technology Attachment

USB: Universal Serial Bus




Figure 36.2: Modern System Architecture

I/O Devices (Chapter 36)

36.2 A Canonical IO Device

Interface: What the
device exposes

Internal: Device
specific hardware and
software

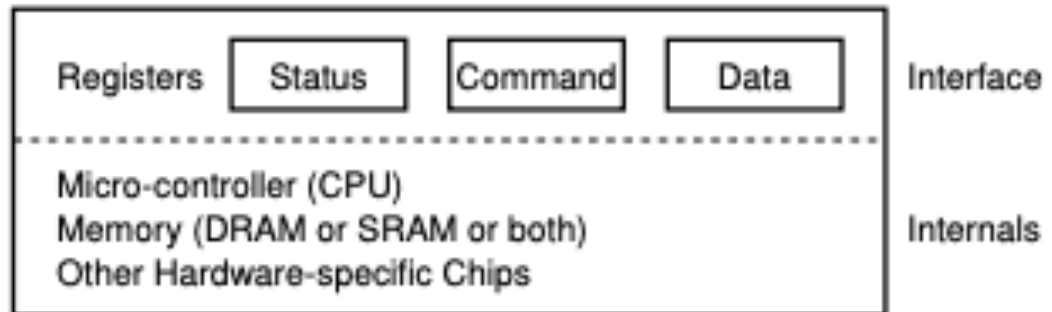


Figure 36.3: A Canonical Device

I/O Devices (Chapter 36)

36.2 A Canonical IO Device

Story:

- A program p1 wants to send sound to a speaker

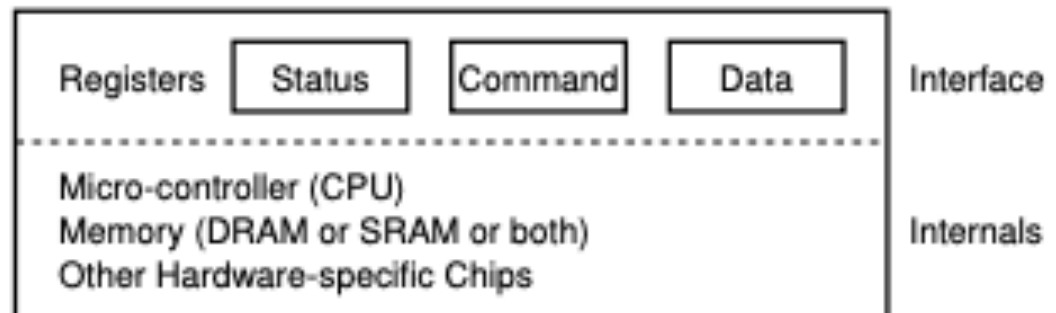


Figure 36.3: A Canonical Device

I/O Devices (Chapter 36)

36.2 A Canonical IO Device

Story:

- A program p1 wants to send sound to a speaker
- p1 issues a system call
- the system call handler traps in the os
- the os calls that procedure that knows how to communicate with the speaker
- How to communicate with the devices comes later 36.6/7

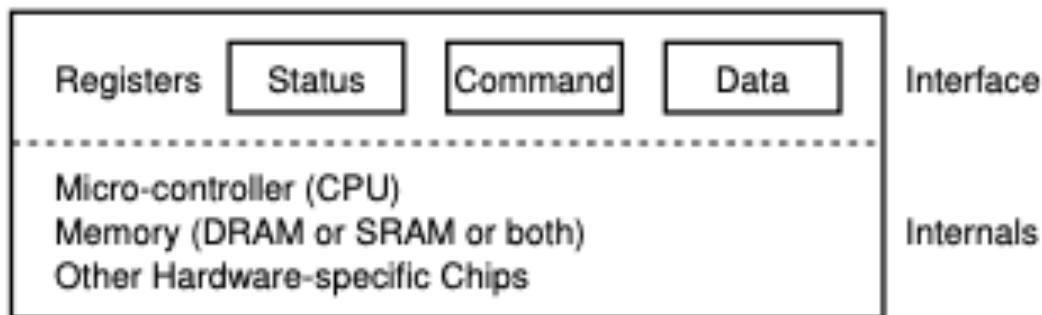


Figure 36.3: A Canonical Device

I/O Devices (Chapter 36)

36.3 A Canonical Protocol

```
While (STATUS == BUSY)
    ; // wait until device is not busy
Write data to DATA register
Write command to COMMAND register
    (starts the device and executes the command)
While (STATUS == BUSY)
    ; // wait until device is done with your request
```

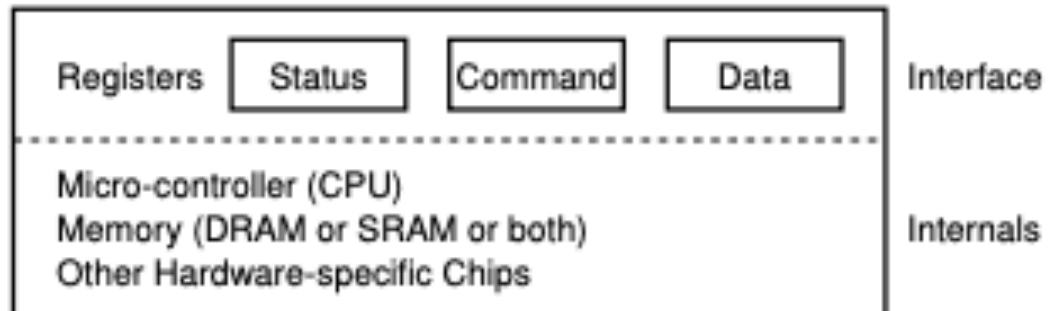


Figure 36.3: A Canonical Device

I/O Devices (Chapter 36)

36.3 A Canonical Protocol

```
While (STATUS == BUSY)
    ; // wait until device is not busy
Write data to DATA register
Write command to COMMAND register
    (starts the device and executes the command)
While (STATUS == BUSY)
    ; // wait until device is done with your request
```

**What is wrong with
this approach ?**




Figure 36.3: A Canonical Device

I/O Devices (Chapter 36)

36.3 A Canonical Protocol

Polling for Ready (CPU)

```
While (STATUS == BUSY)
    ; // wait until device is not busy
Write data to DATA register
Write command to COMMAND register
    (starts the device and executes the command)
While (STATUS == BUSY)
    ; // wait until device is done with your request
```




Figure 36.3: A Canonical Device

I/O Devices (Chapter 36)

36.3 A Canonical Protocol

```
While (STATUS == BUSY)
    ; // wait until device is not busy
Write data to DATA register
Write command to COMMAND register    Writing (CPU)
    (starts the device and executes the command)
While (STATUS == BUSY)
    ; // wait until device is done with your request
```




Figure 36.3: A Canonical Device

I/O Devices (Chapter 36)

36.3 A Canonical Protocol

```
While (STATUS == BUSY)
    ; // wait until device is not busy
Write data to DATA register
Write command to COMMAND register
    (starts the device and executes the command)
While (STATUS == BUSY)    Polling again (CPU)
    ; // wait until device is done with your request
```

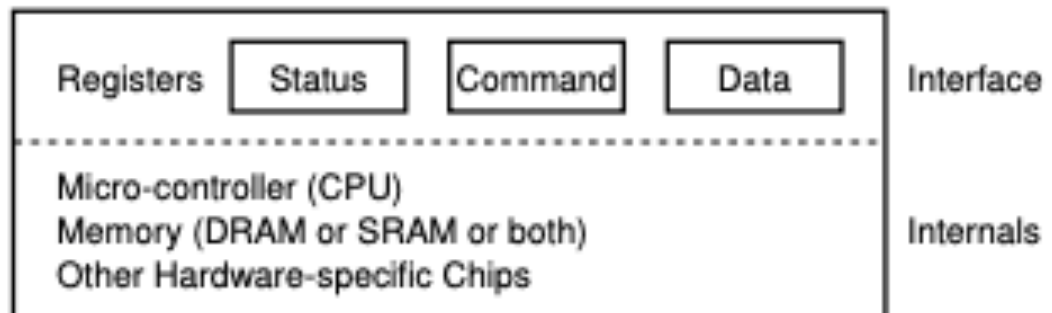


Figure 36.3: A Canonical Device

I/O Devices (Chapter 36)

36.4 Lowering CPU Overhead with Interrupts

Go to sleep and continue later

```
While (STATUS == BUSY)  
    ; // wait until device is not busy  
Write data to DATA register  
Write command to COMMAND register  
    (starts the device and executes the command)  
While (STATUS == BUSY)  
    ; // wait until device is done with your request
```




Figure 36.3: A Canonical Device


I/O Devices (Chapter 36)

36.4 Lowering CPU Overhead with Interrupts

Without interrupt busy spinning

```

While (STATUS == BUSY)
    ; // wait until device is not busy
Write data to DATA register
Write command to COMMAND register
    (starts the device and executes the command)
While (STATUS == BUSY)
    ; // wait until device is done with your request
  
```




I/O Devices (Chapter 36)

36.4 Lowering CPU Overhead with Interrupts

With interrupt another process can use the CPU

```

While (STATUS == BUSY)
    ; // wait until device is not busy
Write data to DATA register
Write command to COMMAND register
    (starts the device and executes the command)
While (STATUS == BUSY)
    ; // wait until device is done with your request
  
```



I/O Devices (Chapter 36)

36.4 Lowering CPU Overhead with Interrupts

```
While (STATUS == BUSY)
    ; // wait until device is not busy
Write data to DATA register
Write command to COMMAND register
    (starts the device and executes the command)
While (STATUS == BUSY)
    ; // wait until device is done with your request
```

However,

if the device is fast, then no need to go to sleep and wait for interrupt

I/O Devices (Chapter 36)

36.4 Lowering CPU Overhead with Interrupts

```
While (STATUS == BUSY)
    ; // wait until device is not busy
Write data to DATA register
Write command to COMMAND register
    (starts the device and executes the command)
While (STATUS == BUSY)
    ; // wait until device is done with your request
```

However,

if the device is fast sometimes, spin sometime and then go to sleep

I/O Devices (Chapter 36)

36.4 Lowering CPU Overhead with Interrupts

```
While (STATUS == BUSY)
    ; // wait until device is not busy
Write data to DATA register
Write command to COMMAND register
    (starts the device and executes the command)
While (STATUS == BUSY)
    ; // wait until device is done with your request
```

Issue,

When many requests (Stream of IO), many interrupts will cause a “livelock” all time goes to processing interrupts

I/O Devices (Chapter 36)

36.4 Lowering CPU Overhead with Interrupts

```
While (STATUS == BUSY)
    ; // wait until device is not busy
Write data to DATA register
Write command to COMMAND register
    (starts the device and executes the command)
While (STATUS == BUSY)
    ; // wait until device is done with your request
```

Issue,

When many requests (Stream of IO), many interrupts will cause a “livelock” all time goes to processing interrupts

I/O Devices (Chapter 36)

36.4 Lowering CPU Overhead with Interrupts

```
While (STATUS == BUSY)
    ; // wait until device is not busy
Write data to DATA register
Write command to COMMAND register
    (starts the device and executes the command)
While (STATUS == BUSY)
    ; // wait until device is done with your request
```

Trick,

The device waits a while before sending interrupt signal to inform that it is done “coalescing” several “done” into one.

I/O Devices (Chapter 36)

36.5 Direct Memory Access (DMA)


```
While (STATUS == BUSY)
    ; // wait until device is not busy
Write data to DATA register
Write command to COMMAND register
    (starts the device and executes the command)
While (STATUS == BUSY)
    ; // wait until device is done with your request
```

Problem,

writing large chunk of data is still a CPU task. Must copy the data from memory to the device one word at the time "not good"

I/O Devices (Chapter 36)

36.5 Direct Memory Access (DMA)




Problem,

writing large chunk of data is still a CPU task. Must copy the data from memory to the device one word at the time "not good"

I/O Devices (Chapter 36)

36.5 Direct Memory Access (DMA)



Solution,

OS (via a system call) tells DMA where the data is in memory, and how much to copy to which device. The device issues an interrupt when done.

I/O Devices (Chapter 36)

36.6 Methods of Device Interaction

Basically how CPU interacts with devices

I/O instructions: uses dedicated commands to communicate with devices

Memory mapped I/O: uses memory operations to communicate with devices

I/O Devices (Chapter 36)

36.6 Methods of Device Interaction

Basically how CPU interacts with devices

Explicit privileged I/O instructions: for example “in”
“out” and port (device)

Memory mapped I/O: device registers are seen as
memory locations. Load (read) or Store (write) to
those locations. The hardware routes that to the
device and not to memory

I/O Devices (Chapter 36)

36.7 The Device Driver

An abstraction make life easier for the OS, but dealing with the device details

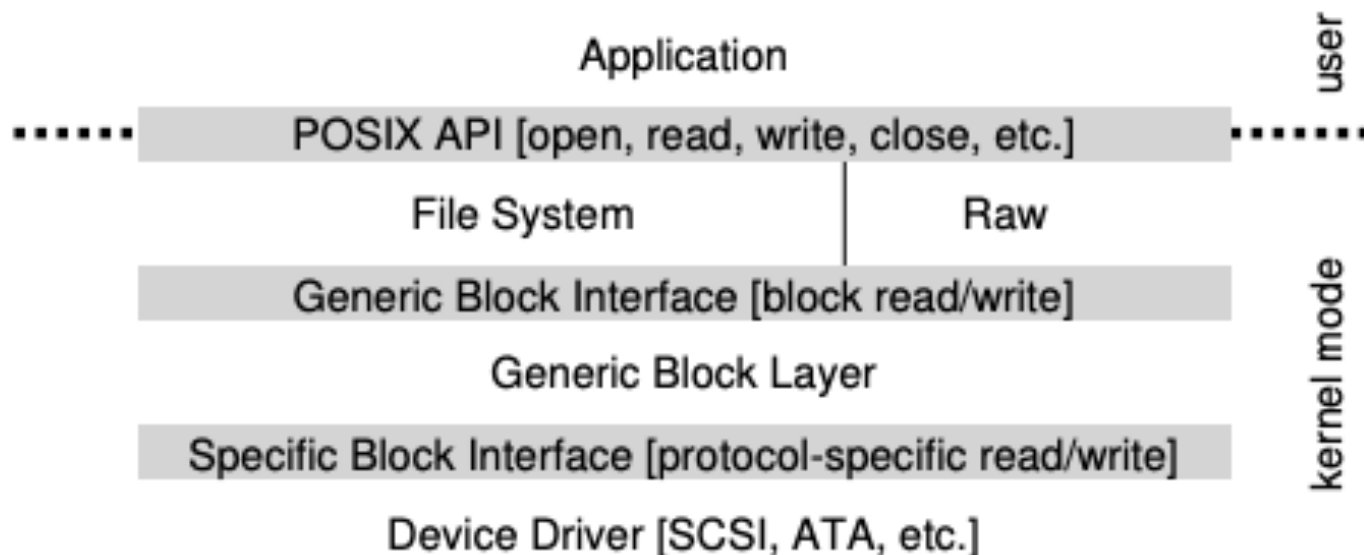


Figure 36.4: The File System Stack

I/O Devices (Chapter 36)

36.10 Summary

- About how OS interacts with a device
- Two techniques for device efficiency: Interrupt and DMA
 - Interrupt makes sense when device is slow
 - DMA makes sense for large data chunks copying
- Two techniques for the cpu to access devices: explicit I/O instructions or memory-mapped I/O
- Device drivers to make it easier for the os to build os in a device neutral fashion.